

Brady's Squad: Human-Led AI Teams

Module 5: Specification-Driven Frameworks | Advanced

Session 16 of 19 | 3 hours | Optional

Agenda

Segment	Time
Coordination trade-offs	7 min
Roles, routing, and memory	20 min
Issues and monitored work	8 min
Prepared team demonstration	12 min
Decision framework and lab handoff	13 min
Lab	2 hours

Coordination has a cost

Use a team when work has separate owners, clear interfaces, and a durable decision to share. An individual agent is better for a small, well-understood change.

More agents create more state, review work, and cost.

Parallelism only helps with real independence

Work can run in parallel when each stream has its own inputs, a clear output, and no hidden dependency on another stream's decision.

Work item	Parallel?	Reason
Review separate test files	Often	Each reviewer can return findings
Choose a shared domain model	No	Every implementation depends on the decision
Update unrelated documentation	Often	Output boundaries are clear

Splitting an ambiguous task across agents spreads the ambiguity instead of removing it.

The coordinator owns the integration risk

A coordinator should not merely hand out tasks. It needs to know:

1. Which decision makes the workstreams compatible.
2. Which files or systems only one role may change.
3. What evidence each role must return.
4. When a conflict should stop the workflow rather than be merged automatically.

This is why the coordinator is a human-led role even when agents do the initial work.

Use a result packet, not a chat transcript

Each workstream should return a small, reviewable artifact:

```
Scope completed:  
Evidence:  
Files or records changed:  
Checks run:  
Open questions:  
Recommendation:
```

The parent can compare packets without reconstructing a long agent conversation.
Missing evidence is a finding, not an invitation to guess.

Coordinator pattern

```
graph TD; A[Issue with acceptance criteria] --> B[Coordinator selects a bounded role]; B --> C[Role writes a proposal or change]; C --> D[Focused checks and human review]; D --> E[Decision record, merge, follow-up, or pause];
```

Issue with acceptance criteria

↓

Coordinator selects a bounded role

↓

Role writes a proposal or change

↓

Focused checks and human review

↓

Decision record, merge, follow-up, or pause

The coordinator routes bounded work. An ambiguous request needs clarification before implementation.

Team boundaries

Role	Responsibility
Lead	Scope, decisions, review gate
Implementer	A bounded product area or file set
Tester	Focused checks and reported gaps
Scribe	Shared decisions and history
Work monitor	Surface ready, approved work

Agents need charters that state inputs, outputs, and limits. Avoid concurrent writes to the same files.

Shared memory

Commit durable decisions: the choice, evidence, owner, and date. Keep temporary drafts separate.

Agents can write individual proposals. A designated owner reviews and merges the shared record. Git history keeps the trail.

Issues and monitored work

An issue is a work contract. It needs acceptance criteria, non-goals, constraints, checks, and a reviewer.

Ready issue → assignment → bounded change → focused checks →
human review → merge or follow-up issue

Treat Ralph as a monitored queue. Stop when approval, ownership, policy, test status, scope, or usage limits are unclear.

When orchestration helps

Situation	Path
One small repair	Individual agent or manual work
Separate work with explicit interfaces	Coordinated team
Overlapping ownership or unclear design	Plan before implementation
Restricted path or unapproved tool	Manual workflow

Demonstration and lab

Inspect a prepared Squad team or initialize one through an approved path. Show the roster, routing rules, charter, and decision record. Assign one narrow issue and review its outcome.

Initialize a team, assign bounded work, use GitHub Issues, and monitor a small queue. One reviewed issue is sufficient.