

# Code Review Workflows

Evidence, human judgment, and policy-approved automation

# Access and cost preflight

---

Use Enterprise Cloud as the governance baseline.

1. Verify current official GitHub documentation and the applicable administrator policy.
2. Confirm repository scope, data classification, and review permissions.
3. Use a non-sensitive pull request with explicit acceptance criteria.
4. For metered work, define an approved threshold, escalation route, and stop guard.

# Review is a human responsibility

---

Automation may raise questions. A reviewer decides whether a change is correct, safe, maintainable, and within scope.

Issue and acceptance criteria



Proposed change and checks



Human review



Accept, request revision, or implement manually

# Review rubric

---

- Correctness against acceptance criteria
- Tests and validation evidence
- Security, dependency, and permission implications
- Scope and maintainability
- Repository conventions and human approval

## Conflict-resolution exercise

---

1. Read both branch intentions.
2. Write the merged behavior and tests before editing.
3. Resolve the conflict manually in the repository sandbox.
4. Run checks and request review.
5. Demonstrate an approved assisted workflow only after checking current documentation and policy.

## No-access fallback

---

Use the intentional defects and conflict scenario for a human-only review. Record findings as pull-request comments, complete the fixes manually, and compare with the solution reference.

# Agenda and time plan

| Time | Topic                       | Trainer move                 |
|------|-----------------------------|------------------------------|
| 0:00 | Review workflow and roles   | Define accountability        |
| 0:08 | Summaries and descriptions  | Separate draft from evidence |
| 0:18 | Manual and automatic review | Show policy controls         |
| 0:28 | End-to-end PR demo          | Apply the rubric             |
| 0:43 | Suggestions and conflicts   | Preserve author intent       |
| 0:51 | Limits and adoption         | Design human gates           |
| 0:58 | Lab handoff and Q&A         | Assign reviewer roles        |

**Transition:** “Copilot can increase review coverage; it cannot own the decision.”

# Copilot across a PR lifecycle

Issue and acceptance criteria



Branch and bounded commits



Draft summary and PR description



Automated checks + Copilot first pass



Human review and author response



Approval, revision, or close

Each stage gives the accountable reviewer evidence.

# Commit messages and PR descriptions

---

Copilot can draft text from the visible change, but authors must verify:

- the user-visible behavior is accurate;
- tests and results are not invented;
- breaking changes and migrations are explicit;
- security and dependency implications are named;
- issue links and acceptance criteria match the work;
- no sensitive details appear in generated prose.

A polished summary can still be wrong.

# A reviewable PR template

---

## **## Problem**

What user-visible problem does this solve?

## **## Change**

List the bounded implementation choices.

## **## Validation**

- [ ] Focused tests
- [ ] Existing checks
- [ ] Manual edge case

## **## Risk and rollback**

Name data, dependency, security, and recovery considerations.

# Requesting Copilot review

---

For a manual request:

1. Confirm the PR is ready for a first pass.
2. Verify current availability, permissions, and repository policy.
3. Request the Copilot review through the supported PR control.
4. Wait for comments and review them as hypotheses.
5. Resolve, reject with rationale, or implement a verified fix.
6. Request human review with the complete evidence.

Do not use Copilot review to bypass required reviewers.

# Automatic review configuration

---

Automatic review can be configured at repository or organization scope when supported.

Before enabling:

- identify eligible repositories and excluded data;
- define when reviews run and who owns triage;
- preserve branch protection and required human approvals;
- monitor noise, missed findings, and review delay;
- document an off switch and escalation owner.

**Trainer cue:** Demonstrate navigation, not fixed UI labels that may change.

# Quality gates

| Gate                   | Evidence                                            |
|------------------------|-----------------------------------------------------|
| Ready for first pass   | Clear description, focused diff, tests available    |
| Ready for human review | Checks pass; AI comments triaged                    |
| Ready for approval     | Acceptance criteria met; risks reviewed             |
| Ready for merge        | Required approvals and branch protections satisfied |

Copilot review comes before human approval. It never replaces it.

## Demo PR: intentional defect

---

```
export function canViewInvoice(user, invoice) {  
  return user.role === "admin" || user.id === invoice.ownerId;  
}
```

The requirement says suspended users must never view invoices.

```
if (user.status === "suspended") return false;  
return user.role === "admin" || user.id === invoice.ownerId;
```

Ask which tests and policy evidence should accompany this suggestion.

## Demo narrative: complete lifecycle

---

1. Open the issue and read acceptance criteria aloud.
2. Inspect the focused branch diff before generated prose.
3. Generate or improve the PR description; correct unsupported claims.
4. Run required checks and request Copilot review.
5. Classify each comment: valid, uncertain, noise, or out of scope.
6. Apply one verified fix and add a regression test.
7. Request human review and record the final decision.

## Reviewing suggested fixes

---

Before applying:

- restate the finding in your own words;
- reproduce the defect or identify the violated requirement;
- inspect the complete patch, not only highlighted lines;
- add or update a test that fails before the fix;
- run broader checks appropriate to the risk;
- retain author ownership of the resulting code.

One-click application saves typing. The reviewer still decides.

## What review may catch

---

- localized correctness issues;
- missing error handling or tests;
- suspicious patterns visible in the diff;
- inconsistent naming or repository conventions;
- simple maintainability concerns.

## What it may miss

---

- unstated product intent;
- cross-system behavior absent from context;
- runtime, operational, legal, or organizational risk;
- subtle authorization and data-boundary flaws;
- deliberately misleading but plausible changes.

# Merge conflict resolution

Base behavior + branch A intent + branch B intent

↓

write merged acceptance criteria

↓

resolve conflict in sandbox

↓

run tests and human review

If an approved cloud-agent workflow is available, compare its proposal with the manual resolution. Do not blindly choose "ours" or "theirs."

# Adoption patterns

---

**Start:** Optional first-pass review on a small set of repositories.

**Measure:** Actionable findings, false positives, review time, rework, and defects that reach users.

**Adapt:** Tune scope, guidance, ownership, and escalation.

**Expand:** Only when quality, security, and developer experience meet team criteria.

Keep a manual path for unavailable features and exceptional repositories.

# Facilitation prompts

---

- Which comment needs domain expertise to evaluate?
- What evidence would turn an uncertain comment into a decision?
- Who resolves review noise?
- Which repositories should not use automatic review?
- How will the team detect over-reliance?

**Transition:** “Now apply the workflow as both author and reviewer.”

## Lab handoff

---

Learners will:

1. follow one pull request from a flawed change brief;
2. classify human and Copilot findings;
3. apply only fixes supported by evidence;
4. resolve the prepared conflict without losing either branch goal;
5. make an explicit final human decision.

**Final deliverable:** The review record, corrected project, conflict evidence, and final decision.

# Remember

---

1. Generated summaries must be verified against the diff.
2. AI review is a first pass before accountable human review.
3. Suggested fixes require reproduction, inspection, and tests.
4. Conflict resolution starts from intended merged behavior.
5. Adoption should be measured and reversible.

# Questions and lab readiness

Which review decision must never be delegated without human evidence?

- Who is the final reviewer?
- Which check is required before requesting review?
- Is the conflict scenario ready?