

Introduction to GitHub Copilot

Define, inspect, test, decide

Session 01 | 3 hours

Agenda

Segment	Time
Proposal - and - evidence model	7 min
How context shapes a suggestion	10 min
Access and data preflight	10 min
Inline suggestion controls	11 min
Prepared utility demonstration	12 min
Follow - up edits	6 min
Fallback and lab handoff	4 min

The durable workflow

```
graph TD; A[Define observable behavior] --> B[Check access, data, and scope]; B --> C[Inspect the proposed change]; C --> D[Run focused checks]; D --> E[Accept, revise, or reject];
```

The developer owns the result.

Proposal is not proof

Proposal	Evidence
Looks plausible	A focused test passes
Uses familiar syntax	The diff stays in scope
Explains its reasoning	A reviewer can reproduce the result
Matches one example	Edge and failure behavior are checked

Fluent output can still be incomplete or wrong.

A practical system model

Task + approved context



Selected Copilot surface



Proposed output



Local checks + human review

Use this as a review model, not a fixed service diagram.

Context changes the proposal

Useful context includes:

- function names and type hints;
- docstrings and examples;
- selected code and nearby files;
- tests and repository instructions.

More context is useful only when it is relevant and approved.

Preflight before use

1. Confirm the account and surface.
2. Confirm the repository and data classification.
3. Limit the task to named files.
4. Name any metered-work owner and stop condition.
5. Prepare the manual route.

If approval is unclear, use the fallback.

Access policy

The live lab requires GitHub Copilot in an approved editor.

The manual fallback uses the same:

- starter code;
- acceptance tests;
- review decisions;
- final verification.

Access changes the tool path, not the learning objective.

Accept, revise, reject

Accept when the proposal meets the contract and fits the repository.

Revise when the direction is useful but a detail is wrong.

Reject when it changes scope, invents behavior, or is hard to verify.

Undo is a normal review action.

Review one suggestion

Before accepting, ask:

- Which requirement does each line support?
- What happens on empty or unusual input?
- Did it add a dependency?
- Did it change a caller or public name?
- Which check will fail if the logic is wrong?

Demo contract

```
def normalize_username(value: str) -> str:  
    """Trim outer whitespace, lowercase the value,  
    and join words with hyphens."""
```

```
assert normalize_username("  River Team  ") == "river-team"
```

The test makes repeated whitespace observable.

Demo: inspect before accepting

Review the proposed implementation against three rules:

1. Outer whitespace disappears.
2. Case becomes lowercase.
3. Each run of internal whitespace becomes one hyphen.

Then run the focused test.

A failing check is useful

This version looks reasonable but mishandles repeated spaces:

```
return value.strip().lower().replace(" ", "-")
```

The failure exposes a hidden assumption while the change is still small.

Predicted follow-up edits still need review

Rename username to handle



Inspect each proposed caller update



Search for the old name



Run the focused test

A predicted edit can miss a reference.

Manual fallback

When live access is unavailable:

1. Implement the function manually.
2. Compare it with one prepared candidate.
3. Mark accept, revise, or reject.
4. Run the same tests.
5. Record the same evidence.

Lab handoff

You will evolve one Python utility module:

1. complete `normalize_username` and `is_palindrome`;
2. decide how apostrophes count in `word_frequency`;
3. add `profile_summary`;
4. rename `username` to `handle`;
5. run the full test suite and review the diff.

Deliverable: completed code, five passing tests, and a decision record.