

Enterprise Governance

A tabletop framework for policy and measurement

Agenda

Segment	Time
Decision owners and evidence	8 min
Access, data, and tool boundary	10 min
Worked evidence map	12 min
Policy decision and checkpoint	12 min
Measurement and rollout gates	10 min
Lab handoff	8 min

Start with evidence

Use Enterprise Cloud as the baseline. Check current GitHub documentation and customer policy before making an access, data, tool, or rollout decision.

A training slide cannot establish feature availability, entitlement, retention, or compliance.

Governance turns uncertainty into decisions

Governance decides what can run, under which limits, and what evidence will change that decision.

Question	Record
Can this use case start?	Scope, owner, data boundary, and approval
Is it safe to continue?	Quality, usage, and review evidence
Can it expand?	The approver and the added boundary
Should it pause?	Trigger, recovery, and next review date

Policy text alone is not a rollout plan. Someone must own the decision.

A boundary has inputs, actions, and outputs

Classify more than the prompt text:

Inputs: repository files, issues, tool results

Actions: read, summarize, write, deploy

Outputs: chat, commit, pull request, external system

Evidence: logs, review, approval, metric

For each item, decide what is allowed, who approves exceptions, and what must stay out of the workflow. This gives the team a working boundary.

Measurements need a decision attached

Signal	Example decision
Review rework rises	Pause expansion and inspect the workflow
Usage reaches a guard	Stop automated runs until the owner reviews
Acceptance rate is stable	Continue the bounded trial
Restricted data appears	Stop, preserve evidence, and escalate

An activity count without a threshold and response tells the team little.

Preflight

Confirm	Why it matters
Sandbox or repository	Keeps the trial bounded
Owner and reviewer	Makes decisions accountable
Data classification	Defines what may enter the workflow
Approved features and tools	Prevents an unsupported path
Meter and stop guard	Limits metered work
Manual fallback	Keeps learning moving when access is blocked

Policy decision sequence

Evidence map → bounded trial → observed review evidence →
continue, expand, gather more evidence, or pause

Unknown is a decision state. Record the owner and needed evidence. **Unknown does not mean approved.**

Worked scenario

The fictional team proposes a two-week trial:

Field	Decision
Repository	training-review-sandbox
Workflow	Review one synthetic pull request at a time
Data	Synthetic task records only
Tools	Approved repository assistant; no external MCP server
Excluded	Production, deployment, customer data, live policy changes
Reviewer	Engineering lead
Stop	Restricted data, failed boundary check, or 20 automated runs

The scenario gives the class fixed facts. Current product and policy claims still need live evidence from the proper owner.

Evidence-map checkpoint

Record each claim as:

- **confirmed in scenario;**
- **live evidence pending;**
- **out of scope.**

The learner may continue only when the repository, workflow, data class, tool boundary, reviewer, meter, and fallback have an owner and status.

Policy decision checkpoint

Reference decision: Proceed in the repository sandbox only.

The trial may start with synthetic data and human review. Expansion remains blocked until the administrator and data owner confirm current policy evidence.

If a learner cannot state the approved and blocked scope in one sentence, the decision record is not ready.

Data and tool boundaries

Classify the repository and inputs before using an assistant. Keep restricted content out of prompts, issues, comments, and training tools.

Review MCP servers, approved package sources, and external tools the same way: owner, purpose, permitted data, permissions, logging, and fallback.

Content exclusion is one control. Check current surface support, and keep restricted data out of the workflow.

Measurement guardrail

Activity is not value. Compare bounded work with acceptance criteria, human review, safety evidence, and a known baseline.

Guardrail item	Required decision
Exercise guard	Trainer - owned counter for the scenario
Window	When observations are collected
Threshold	When an alert fires
Escalation	Who responds
Stop	When automated work pauses
Fallback	Manual work while evidence is pending

The exercise run counter is a stop guard. It is not a product usage or billing metric.

Worked measurement rule

Item	Fictional value
Window	Two weeks or 10 reviewed tasks
Quality	At least 9 tasks meet acceptance criteria
Safety	Zero restricted-data events
Rework	No more than 3 tasks need material rewrite
Exercise guard	Stop at 20 automated runs
Response	Pause, preserve evidence, and use manual review

These values support the exercise. They are not product limits or commercial guidance.

Rollout gates

Gate	Evidence
Start	Approved policy, data boundary, tools, reviewer, and fallback
Continue	Quality, safety, review, and meter evidence
Expand	Explicit approval from required owners
Pause	Unresolved policy, data, quality, or measurement concern

Lab handoff

Use the fixed `training-review-sandbox` scenario.

1. Complete the evidence map and pass Checkpoint 1.
2. Record the sandbox-only policy decision and pass Checkpoint 2.
3. Define the measurement guardrail.
4. Complete start, continue, expand, and pause gates.

Submit four completed artifacts. Compare them with `lab/solution/` after the final decision, not before.